

Towards Verifying the Bitcoin-S Library

Ramon Boss, Kai Brännler, and Anna Doukmak

Bern University of Applied Sciences, CH-2501 Biel, Switzerland
ramon.boss@outlook.com, kai.brueennler@bfh.ch,
anna.doukmak@gmail.com

Abstract. We try to verify properties of the bitcoin-s library, a Scala implementation of parts of the Bitcoin protocol. We use the Stainless verifier which supports programs in subset of Scala called the *Pure Scala Fragment*. We first try to verify the property that regular transactions do not create new money. It turns out that there is too much code involved that lies outside of the supported fragment to make this feasible. However, in the process we uncover and fix a bug in bitcoin-s. We then turn to a much simpler (and less interesting) property: that adding zero satoshis to a given amount of satoshis yields the given amount of satoshis. Here as well a significant part of the relevant code lies outside of the supported fragment. However, after a series of equivalent transformations we arrive at code that we successfully verify.

Keywords: Bitcoin · Scala · Bitcoin-S · Stainless.

1 Introduction

For software handling cryptocurrency, correctness is clearly crucial. However, even in very well-tested software such as Bitcoin Core, serious bugs occur. The most recent example is the bug found in September 2018 [7] which essentially allowed to arbitrarily create new coins. Such software is thus a worthwhile target for formal verification. In this work, we set out to verify properties of the bitcoin-s library with the Stainless verifier.

The Bitcoin-S Library. The bitcoin-s library is an implementation of parts of the Bitcoin protocol in Scala [9,10]. In particular, it allows to serialize, deserialize, sign and validate transactions. The library uses immutable data structures and algebraic data types but is not written with formal verification in mind. According to the website, the library is used in production, handling significant amounts of cryptocurrency each day [9].

The Stainless Verifier. Stainless is the successor of the Leon verifier [2,11,1] and is developed at EPF Lausanne [5]. It is intended to be used by programmers without training in formal verification and thus allows to write specifications in Scala and focusses on counterexample finding in addition to proving correctness.

The example in Figure 1 from the Stainless documentation [8] demonstrates this. Notice how a precondition is specified using the function *require* and a postcondition using *ensuring*.

```

1  def factorial(n: Int): Int = {
2      require(n >= 0)
3      if (n == 0) {
4          1
5      } else {
6          n * factorial(n - 1)
7      }
8  } ensuring(res => res >= 0)

```

Fig. 1. Factorial program with specification

Our program happens not to satisfy the specification. An overflow in the 32-bit Int leads to a negative result for the input 17, as Stainless reports in Figure 2. Changing the type from Int to BigInt will result in a successful verification.

```

[Warning ] The Z3 native interface is not available. Falling back onto smt-z3.
[ Info ] - Checking cache: 'postcondition' VC for factorial @10:3...
[ Info ] - Checking cache: 'precond. (call factorial(n - 1))' VC for factorial @15:11...
[ Info ] - Checking cache: 'postcondition' VC for factorial @10:3...
[ Info ] Cache miss: 'postcondition' VC for factorial @10:3...
[ Info ] Cache hit: 'precond. (call factorial(n - 1))' VC for factorial @15:11...
[ Info ] Cache hit: 'postcondition' VC for factorial @10:3...
[ Info ] - Now solving 'postcondition' VC for factorial @10:3...
[ Info ] - Result for 'postcondition' VC for factorial @10:3:
[Warning ] => INVALID
[Warning ] Found counter-example:
[Warning ] n: Int -> 17
[ Info ]
[ Info ] stainless summary
[ Info ]
[ Info ] factorial postcondition          valid from cache      src/TestFactorial.scala:10:3  1.055
[ Info ] factorial postcondition          invalid              U:smt-z3 src/TestFactorial.scala:10:3  7.861
[ Info ] factorial precondition (call factorial(n - 1)) valid from cache      src/TestFactorial.scala:15:11  1.054
[ Info ]
[ Info ] total: 3   valid: 2   (2 from cache) invalid: 1   unknown: 0   time: 9.970
[ Info ]
[ Info ] Shutting down executor service.

```

Fig. 2. Output of Stainless verification for calculating factorial of Int number

The Pure Scala Fragment. The Scala fragment supported by Stainless is described in the Stainless documentation [8] in the section [Pure Scala](#).

It comprises algebraic data types in the form of abstract classes, case classes and case objects, objects for grouping classes and functions, boolean expressions with short-circuit interpretation, generics with invariant type parameters, default values of function parameters, pattern matching, local and anonymous classes and more.

In addition to Pure Scala Stainless also supports some imperative features, such as using a (mutable) variable in a local scope of a function and while loops. They turn out not to be relevant for the current work.

What will turn out to be more relevant are the following Scala features which Stainless does not support, such as: (concrete) class definitions, inheritance by objects, abstract type members, variance annotations and private inner classes.

In addition, Stainless has its own library of some core data types and functions which need to be mapped correctly to functions inside of the SMT solver that Stainless ultimately relies on. Those data types in general do not have all the methods of the Scala data types. For example, the `BigInt` type in Scala has a methods for bitwise operations while the `BigInt` type in Stainless does not.

Outline and Properties to Verify. In the next section we try to verify the property that a regular (non-coinbase) transaction cannot generate new coins. We call it the *no-inflation property*. Trying to verify it, we uncover and fix a bug in the bitcoin-s library. We then find that there is too much code involved that lies outside of the supported fragment to currently make this verification feasible. Instead, we turn to a simpler property to verify. The simplest possible property we can think of is the fact that adding zero satoshis to a given amount of satoshis yields the given amount of satoshis. We call it the *addition-with-zero property* and we try to verify it in Section 3. Here as well we see that a significant part of the code lies outside of the supported fragment. However, after a series of equivalent transformations we arrive at code that we successfully verify.

2 The No-Inflation Property

A crucial function for the verification of the no-inflation property is the `checkTransaction` function, shown in Figure 5. To better understand it, we first see how to create a transaction.

Creating a Transaction

To create a transaction, we first need some coins – an unspent transaction output. We could load an actual unspent transaction output from the bitcoin network, but we create one manually in order to see this process. So we first create an (invalid) transaction with one output in Figure 3.

We first create a keypair, then a lock script with its public key, then the amount of satoshis, then a transaction output (utxo) for that amount and locked with that script. Finally we create the actual transaction with that output and no inputs. Of course, that is not a valid transaction, because it creates coins out of nothing. In particular, `checkTransaction(prevTx)` is false.

Now that we have a transaction output, we create a new transaction to spend it.

First, we need some out points. They point to outputs of previous transactions. We use the index zero, because the previous transaction has only one output that becomes the first index zero. If there were two previous outputs, the second output would become the index 1 and so on.

This utxos are the inputs of our transaction. Second, we need destinations to spend the bitcoins to. For the sake of convenience we create only one. We spend 5000 satoshis to the newly created random public key. Finally, we define the fee rate in satoshis per one byte transaction size as well as some bitcoin network

```

1  val privKey = ECPrivateKey.freshPrivateKey
2  val creditingSPK = P2PKHScriptPubKey(pubKey = privKey.publicKey)
3
4  val amount = Satoshi(Int64(10000))
5
6  val utxo = TransactionOutput(currencyUnit = amount, scriptPubKey =
    creditingSPK)
7
8  val prevTx = BaseTransaction(
9    version = Int32.one,
10   inputs = List.empty,
11   outputs = List(utxo),
12   lockTime = UInt32.zero
13 )

```

Fig. 3. Creating a transaction output to spend

parameters. The bitcoin network parameters are not important, so we use some static values normally used when testing.

Now lets build the transaction with those data. Line one to seven creates a transaction builder which is then signed on line ten. We can now use our transaction object on line twelve. For example, after calling *hex* on it, we can send the returned string to the bitcoin network.

Validating a Transaction

Bitcoin-S offers a function called *checkTransaction* located in the *ScriptInterpreter* object.

We can pass a transaction and it returns a Boolean indicating whether the transaction is valid or not. So for example when we pass the transaction we built before the returned value would be true, because it's a valid transaction. It might not be accepted by the bitcoin network but for a transaction on its own it's valid. We can not check context with it, because we can only pass one transaction.

There are several checks in *checkTransaction*. For example, it checks if there is either no input or no output. In this case we get false.

The relevant part for the bug we found:

```

1  val prevOutputTxIds = transaction.inputs.map(_.previousOutput.txId)
2  val noDuplicateInputs = prevOutputTxIds.distinct.size ==
    prevOutputTxIds.size

```

It gathers all transaction ids referenced by the out points. When we call *distinct* on the returned list, we get a list with duplicate removed. If the size of the new list is the same as the size of the old, we know that there was no duplicate transaction id, because, as said, *distinct* removes the duplicates.

Trying Stainless on the entire project.

– result sbt: no output

```

1  val outPoint = TransactionOutPoint(prevTx.txId, UInt32.zero)
2
3  val utxoSpendingInfo = BitcoinUTXOSpendingInfo(
4    outPoint = outPoint,
5    output = utxo,
6    signers = List(privKey),
7    redeemScriptOpt = None,
8    scriptWitnessOpt = None,
9    hashType = HashType.sigHashAll
10 )
11
12 val utxos = List(utxoSpendingInfo)
13
14 val destinationAmount = Satoshis(Int64(5000))
15
16 val destinationSPK = P2PKHScriptPubKey(pubKey = ECPrivateKey.
17   freshPrivateKey.publicKey)
18
19 val destinations = List(
20   TransactionOutput(currencyUnit = destinationAmount, scriptPubKey
21     = destinationSPK)
22 )
23
24 val feeRate = SatoshisPerByte(Satoshis.one)
25
26 val networkParams = RegTest // some static values for testing
27
28 val txBuilderF: Future[BitcoinTxBuilder] = BitcoinTxBuilder(
29   destinations = destinations, // where to send the money
30   utxos = utxos,              // unspent transaction outputs
31   feeRate = feeRate,          // fee rate per byte
32   changeSPK = creditingSPK,   // where to send the change
33   network = networkParams     // bitcoin network information
34 )
35
36 val txF: Future[Transaction] = txBuilderF.flatMap(_.sign)
37
38 val tx: Transaction = Await.result(signedTxF, 1 second)

```

Fig. 4. Creating a transaction

```

1  /**
2   * Checks the validity of a transaction in accordance to bitcoin
3   * core's CheckTransaction function
4   * https://github.com/bitcoin/bitcoin/blob/
5   * f7a21dae5dbf71d5bc00485215e84e6f2b309d0a/src/main.cpp#L939.
6   */
7  def checkTransaction(transaction: Transaction): Boolean = {
8    val inputOutputsNotZero =
9      !(transaction.inputs.isEmpty || transaction.outputs.isEmpty)
10    val txNotLargerThanBlock = transaction.bytes.size < Consensus.
11      maxBlockSize
12    val outputsSpendValidAmountsOfMoney = !transaction.outputs.exists(o
13      =>
14        o.value < CurrencyUnits.zero || o.value > Consensus.maxMoney)
15    val outputValues = transaction.outputs.map(_.value)
16    val totalSpentByOutputs: CurrencyUnit =
17      outputValues.fold(CurrencyUnits.zero)(_ + _)
18    val allOutputsValidMoneyRange = validMoneyRange(totalSpentByOutputs
19      )
20    val prevOutputTxIds = transaction.inputs.map(_.previousOutput.txId)
21    val noDuplicateInputs = prevOutputTxIds.distinct.size ==
22      prevOutputTxIds.size
23
24    val isValidScriptSigForCoinbaseTx = transaction.isCoinbase match {
25      case true =>
26        transaction.inputs.head.scriptSignature.asmBytes.size >= 2 &&
27        transaction.inputs.head.scriptSignature.asmBytes.size <= 100
28      case false =>
29        //since this is not a coinbase tx we cannot have any empty
30        //previous outs inside of inputs
31        !transaction.inputs.exists(_.previousOutput ==
32          EmptyTransactionOutPoint)
33    }
34    inputOutputsNotZero && txNotLargerThanBlock &&
35      outputsSpendValidAmountsOfMoney && noDuplicateInputs &&
36      allOutputsValidMoneyRange && noDuplicateInputs &&
37      isValidScriptSigForCoinbaseTx
38  }

```

Fig. 5. The checkTransaction function in the ScriptInterpreter object

– result jar:

In order to verify a project, Stainless must be integrated into it. We can integrate it in an sbt project adding the Stainless Plugin and the required resolver to pugins.sbt. Another option to use Stainless is to import its libraries in a program code and verify a program from command line using the pre-packaged Stainless JAR file or using the Stainless script built from source. Trying to integrate Stainless in Bitcoin-S caused a lot of troubles, mainly because of version conflicts. For more details see chapter ??.

After integrating the Stainless plugin in the Bitcoin-S sbt project, there were many errors because of the different sbt versions. Some errors are described in the section 4.2. It takes too much time to fix them all so it should be easier to extract the classes needed for the `checkTransaction` function.

Putting aside the No-Inflation Property. Naively trying Stainless on the entire bitcoin-s codebase results in either no output (with sbt) or many errors (with jar). So we extract the relevant code to only verify that. However, the extracted code has more than 1500 lines and liberally uses Scala features outside of the supported fragment. We tried to transform the code into the supported fragment, but realize that a better approach is to first verify a simpler property with less code involved and then turn back to the no-inflation property with more experience. So we turn to the addition-with-zero property in the next section.

explain

Fixing a Bug in Bitcoin-S

We can see that there is a bug in the `checkTransaction` function from before, recognized and fixed through this work.

Here is the relevant code of `checkTransaction` again:

```
1  val prevOutputTxIds = transaction.inputs.map(_.previousOutput.txId)
2  val noDuplicateInputs = prevOutputTxIds.distinct.size ==
    prevOutputTxIds.size
```

What happens if we have two `TransactionOutPoints` (previousOutputs) with a different index but referencing the same Transaction ID (txId)?

According to the Bitcoin protocol this is possible. A transaction can have multiple outputs that should be referenceable by the next transaction. So this is clearly a bug.

What should not be possible is a transaction referencing the same output twice. This bug occurred in Bitcoin Core known as CVE-2018-17144 which was patched on September 18, 2018. [7]

Here, Bitcoin-S did a bit too much and marked all transaction as invalid, if they referenced the same transaction twice. The fix is, to check on `TransactionOutPoint` instead of `TransactionOutPoint.txId`, because `TransactionOutPoint` contains the txId as well as the output index it references. So in pseudo code, we check on the tuple (tx, index) instead of (tx). The fixed code:

```
1  val prevOutputs = transaction.inputs.map(_.previousOutput)
2  val noDuplicateInputs = prevOutputs.distinct.size == prevOutputs.
    size
```

Since `TransactionOutPoint` is a case class and Scala has a built in `==` for case classes there is no need to implement `TransactionOutPoint.==`.

This was fixed in [pull request number 435](#) on GitHub at April 23, 2019, through this work along with a unit test to prevent this bug from appearing again in the future.

3 The Addition-with-Zero Property

In Bitcoin-S there is a class `Satoshis` representing an amount of bitcoins. We look at the verification of the addition of `Satoshis` with zero `Satoshis`. This operation should result in the same amount of `Satoshis`. Let's call it the `??`.

Using Stainless, we see the successful verification of this property. But the process of the verification with the tool requires many changes in the code, so that Stainless can accept it. We look at all needed modifications in chapter 3.

After realizing that it would consume too much time to rewrite the Bitcoin-S code and even the extracted part with `checkTransaction`, the smallest unit in Bitcoin-S-Core that is worthwhile to verify was extracted. This could be the addition of two `CurrencyUnits`. To make it even easier, the addition of `CurrencyUnits` with zero. `CurrencyUnits` is an abstract class in Bitcoin-S, representing currencies like `Satoshis`.

Extracting the relevant Code

The relevant code is in two files: `CurrencyUnits.scala` and `NumberType.scala`. reference: `code/addition/src/main/scala/addition/original`

From those files we removed all code that is not needed for our verification. For example, we removed all number types except for `Int64` (so `Int32`, `UInt64`, etc.) because we do not use them. We also removed the superclasses `Factory` and `NetworkElement` of `CurrencyUnit` and `Number`, respectively because the inherited members are not used by the relevant code. Also, we removed all binary operations on `Number` that are not used.

- removed extending `Factory`, `NetworkElement` and `BasicArithmetic`. This includes some hex/byte conversion eg `fromHex`, `hex`, `bytes`, `fromBytes`, ... Just interfaces never referenced [description in section #the-basics \(cannot add link with hashtag\)](#) `NetworkElement` class `Factory` class

- removed `Bitcoins` class
- removed subtraction and multiplication, binary operations `«`, `»`, etc, comparison operator `<=`, `>=`, etc but not `==` and `!=`
- removed `toBigDecimal`
- removed object `CurrencyUnits` containing some variables to transform `satoshis` to `btc` (not used)

The code we use for the following sections is in `reduced` folder.

Here we can see the extracted code needed for the addition of `CurrencyUnits`:

```
1 package addition.reduced.number
2
3 /**
```



```

4   * This abstract class is meant to represent a signed and unsigned
      number in C
5   * This is useful for dealing with codebases/protocols that rely on
      C's
6   * unsigned integer types
7   */
8   sealed abstract class Number[T <: Number[T]] {
9     type A = BigInt
10
11    /** The underlying scala number used to hold the number */
12    protected def underlying: A
13
14    def toLong: Long = toBigInt.bigInteger.longValueExact()
15    def toBigInt: BigInt = underlying
16
17    /**
18     * This is used to determine the valid amount of bytes in a number
19     * for instance a UInt8 has an andMask of 0xff
20     * a UInt32 has an andMask of 0xffffffff
21     */
22    def andMask: BigInt
23
24    /** Factory function to create the underlying T, for instance a
25        UInt32 */
26    def apply: A => T
27
28    def +(num: T): T = apply(checkResult(underlying + num.underlying))
29
30    /**
31     * Checks if the given result is within the range
32     * of this number type
33     */
34    private def checkResult(result: BigInt): A = {
35      require((result & andMask) == result,
36        "Result was out of bounds, got: " + result)
37      result
38    }
39
40    /**
41     * Represents a signed number in our number system
42     * Instances of this is [[Int64]]
43     */
44    sealed abstract class SignedNumber[T <: Number[T]] extends Number[T]
45
46    /**
47     * Represents a int64_t in C
48     */
49    sealed abstract class Int64 extends SignedNumber[Int64] {
50      override def apply: A => Int64 = Int64(_)

```

```

51   override def andMask = 0xffffffffffffffffL
52 }
53
54 /**
55  * Represents various numbers that should be implemented
56  * inside of any companion object for a number
57  */
58 trait BaseNumbers[T] {
59   def zero: T
60   def one: T
61   def min: T
62   def max: T
63 }
64
65 object Int64 extends BaseNumbers[Int64] {
66   private case class Int64Impl(underlying: BigInt) extends Int64 {
67     require(underlying >= -9223372036854775808L,
68            "Number was too small for a int64, got: " + underlying)
69     require(underlying <= 9223372036854775807L,
70            "Number was too big for a int64, got: " + underlying)
71   }
72
73   lazy val zero = Int64(0)
74   lazy val one = Int64(1)
75
76   lazy val min = Int64(-9223372036854775808L)
77   lazy val max = Int64(9223372036854775807L)
78
79   def apply(long: Long): Int64 = Int64(BigInt(long))
80
81   def apply(bigInt: BigInt): Int64 = Int64Impl(bigInt)
82 }

```



```

1 package addition.reduced.currency
2
3 import addition.reduced.number.{BaseNumbers, Int64}
4
5 sealed abstract class CurrencyUnit {
6   type A
7
8   def satoshis: Satoshi
9
10  def !=(c: CurrencyUnit): Boolean = !(this == c)
11
12  def ==(c: CurrencyUnit): Boolean = satoshis == c.satoshis
13
14  def +(c: CurrencyUnit): CurrencyUnit = {
15    Satoshi(satoshis.underlying + c.satoshis.underlying)
16  }
17

```

```

18   protected def underlying: A
19 }
20
21 sealed abstract class Satoshi extends CurrencyUnit {
22   override type A = Int64
23
24   override def satoshis: Satoshi = this
25
26   def toBigInt: BigInt = BigInt(toLong)
27
28   def toLong: Long = underlying.toLong
29
30   def ==(satoshis: Satoshi): Boolean = underlying == satoshis.
    underlying
31 }
32
33 object Satoshi extends BaseNumbers[Satoshi] {
34
35   val min = Satoshi(Int64.min)
36   val max = Satoshi(Int64.max)
37   val zero = Satoshi(Int64.zero)
38   val one = Satoshi(Int64.one)
39
40   def apply(int64: Int64): Satoshi = SatoshiImpl(int64)
41
42   private case class SatoshiImpl(underlying: Int64) extends Satoshi
43 }

```

The additions' signature looks like this:

```
1 +(c: CurrencyUnit): CurrencyUnit
```

When we run Stainless on this code (without any properties to prove), it throws the following errors:

describe errors: no support for abstract types, unsupported arguments for the BigInt constructor, unsupported inheritance for objects.

Transforming the Code.

kai todo

The Specification.

As specified, our verification must only support addition with zero. So we restrict the parameter to be zero in the precondition.

```
1 require(c.satoshis == Satoshi.zero)
```

We ensure the result is the same value as *this* in the postcondition.

```
1 ensuring(res => res.satoshis == this.satoshis)
```

We can use equals (==) directly on Satoshi, because it is a case class.

Before:

```
1 sealed abstract class CurrencyUnit {
```

```

2   def satoshis: Satoshi
3
4   def !=(c: CurrencyUnit): Boolean = !(this == c)
5
6   def ==(c: CurrencyUnit): Boolean = satoshis == c.satoshis
7
8   def +(c: CurrencyUnit): CurrencyUnit = {
9       Satoshi(satoshis.underlying + c.satoshis.underlying)
10  }
11
12  protected def underlying: Int64
13  }

```

The addition will now look like this:

```

1 sealed abstract class CurrencyUnit {
2     def satoshis: Satoshi
3
4     def !=(c: CurrencyUnit): Boolean = !(this == c)
5
6     def ==(c: CurrencyUnit): Boolean = satoshis == c.satoshis
7
8     override def +(c: CurrencyUnit): CurrencyUnit = {
9         require(c.satoshis == Satoshi.zero)
10        Satoshi(satoshis.underlying + c.satoshis.underlying)
11    } ensuring(res => res.satoshis == this.satoshis)
12
13    protected def underlying: Int64
14 }

```

That's all we need to verify our addition.

Result

Finally, everything is green and correctly verified.

```
[ Info]
[ Info]
[ Info] stainless summary
[ Info]
[ Info] + precondition
[ Info] + precond. (call +(underlying(satoshi(this)), underlying ...))
[ Info] + precond. (call checkResult(this, underlying(this) + ...))
[ Info] + apply
[ Info] + apply precond. (call apply(Int64(), bigInt))
[ Info] + apply add invariant
[ Info] + max precond. (call apply(Int64(), 9223372036854775807))
[ Info] + min precond. (call apply(Int64(), -9223372036854775808))
[ Info] + one precond. (call apply(Int64(), 1))
[ Info] + zero precond. (call apply(Int64(), 0))
[ Info]
[ Info]
[ Info] total: 10 valid: 10 (0 from cache) invalid: 0 unknown: 0 time: 4.792
[ Info]
```

Fig. 6. Output of Stainless verification for addition with 0 of Bitcoin-S-Cores CurrencyUnit

The verified code.

```
1 package addition.modified.number
2
```

```

3  /**
4   * This abstract class is meant to represent a signed and unsigned
      number in C
5   * This is useful for dealing with codebases/protocols that rely on
      C's
6   * unsigned integer types
7   */
8  sealed abstract class Number {
9    /** The underlying scala number used to hold the number */
10   protected def underlying: BigInt
11
12   def toBigInt: BigInt = underlying
13
14   /** Factory function to create the underlying T, for instance a
      UInt32 */
15   def apply: BigInt => Int64
16
17   def +(num: Int64): Int64 = apply(checkResult(underlying + num.
      underlying))
18
19   /**
20    * Checks if the given result is within the range
21    * of this number type
22    */
23   private def checkResult(result: BigInt): BigInt = {
24     require(
25       result <= BigInt("9223372036854775807")
26       && result >= BigInt("-9223372036854775808"))
27     result
28   }
29 }
30
31 /**
32  * Represents a signed number in our number system
33  * Instances of this is [[Int64]]
34  */
35 sealed abstract class SignedNumber extends Number
36
37 /**
38  * Represents a int64_t in C
39  */
40 sealed abstract class Int64 extends SignedNumber {
41   override def apply: BigInt => Int64 = Int64(_)
42 }
43
44 /**
45  * Represents various numbers that should be implemented
46  * inside of any companion object for a number
47  */
48 trait BaseNumbers[T] {

```

```

49   def zero: T
50   def one: T
51   def min: T
52   def max: T
53 }
54
55 case object Int64 extends BaseNumbers[Int64] {
56   lazy val zero = Int64(0)
57   lazy val one = Int64(1)
58
59   lazy val min = Int64(BigInt("-9223372036854775808"))
60   lazy val max = Int64(BigInt("9223372036854775807"))
61
62   def apply(bigInt: BigInt): Int64 = Int64Impl(bigInt)
63 }
64
65 private case class Int64Impl(underlying: BigInt) extends Int64 {
66   require(underlying >= BigInt("-9223372036854775808"))
67   require(underlying <= BigInt("9223372036854775807"))
68 }

```



```

1 package addition.modified.currency
2
3 import addition.modified.number.{BaseNumbers, Int64}
4 import stainless.lang._
5
6 sealed abstract class CurrencyUnit {
7   def satoshis: Satoshi
8
9   def !=(c: CurrencyUnit): Boolean = !(this == c)
10
11   def ==(c: CurrencyUnit): Boolean = satoshis == c.satoshis
12
13   def +(c: CurrencyUnit): CurrencyUnit = {
14     Satoshi(satoshis.underlying + c.satoshis.underlying)
15   } ensuring(res =>
16     (c.satoshis == Satoshi.zero) ==>
17     (res.satoshis == this.satoshis))
18
19   protected def underlying: Int64
20
21 sealed abstract class Satoshi extends CurrencyUnit {
22   override def satoshis: Satoshi = this
23
24   def toBigInt: BigInt = underlying.toBigInt
25
26   def ==(satoshis: Satoshi): Boolean = underlying == satoshis.
    underlying
27 }
28

```

```

29 case object Satoshi extends BaseNumbers[Satoshi] {
30   val min = Satoshi(Int64.min)
31   val max = Satoshi(Int64.max)
32   val zero = Satoshi(Int64.zero)
33   val one = Satoshi(Int64.one)
34
35   def apply(int64: Int64): Satoshi = SatoshiImpl(int64)
36 }
37
38 private case class SatoshiImpl(underlying: Int64) extends Satoshi

```

4 Conclusion

Because of the limitations of the verification tool, we could only verify a rewritten version of the original Bitcoin-S code. So we can not guarantee the correctness of the addition of Satoshi with zero in Bitcoin-S. Not all changes we made were as trivial as the replacement of objects with case objects. For these non-trivial changes, as seen for example the bound check in section ??, we cannot say whether they are equivalent to the original implementation or not.

So code should be written specically with formal verification in mind, in order to successfully verify it. Otherwise, it needs a lot of changes in the software because verification is mathematical and the current software is written mostly in object-oriented style. Software written in the functional paradigm would be much easier to reason about.

Thus, either Stainless must find ways to translate more of built-in object-oriented patterns of Scala to their verification tool or developers must invest more in functional programming.

Also, we found that trying to verify code reveals bugs as shown in section ??. Finally, our work led to some feedback to the Stainless developers to improve the tool.

conclusions: what's future work? how to change bitcoin-s? how to extend stainless?

References

1. Blanc, R., Kuncak, V.: Sound reasoning about integral data types with a reusable SMT solver interface. In: Haller and Miller [6], pp. 35–40. <https://doi.org/10.1145/2774975.2774980>
2. Blanc, R., Kuncak, V., Kneuss, E., Suter, P.: An overview of the leon verification system: verification by translation to recursive functions. In: Proceedings of the 4th Workshop on Scala, SCALA@ECOOP 2013, Montpellier, France, July 2, 2013. pp. 1:1–1:10. ACM (2013). <https://doi.org/10.1145/2489837.2489838>
3. Boss, R.: Issue 519: Unknown type parameter type t in self referencing generic, <https://github.com/epfl-lara/stainless/issues/519>, accessed 2019-06-27
4. developers, T.S.: Pull request 470: Type aliases, type members, and dependent function types, <https://github.com/epfl-lara/stainless/pull/470>, accessed 2019-06-27

5. developers, T.S.: The stainless repository, <https://github.com/epfl-lara/stainless>, accessed 2019-06-19
6. Haller, P., Miller, H. (eds.): Proceedings of the 6th ACM SIGPLAN Symposium on Scala, Scala@PLDI 2015, Portland, OR, USA, June 15-17, 2015. ACM (2015)
7. Song, J.: Bitcoin Core Bug CVE-2018-17144: An Analysis, <https://hackernoon.com/bitcoin-core-bug-cve-2018-17144-an-analysis-f80d9d373362>, accessed 2019-06-20
8. Stainless documentation: <https://epfl-lara.github.io/stainless/>, accessed 2019-06-19
9. Suredbits & the bitcoin-s developers: The bitcoin-s website, <https://bitcoin-s.org>, accessed 2019-06-19
10. The bitcoin-s developers: The bitcoin-s repository, <https://github.com/bitcoin-s>, accessed 2019-06-19
11. Voirol, N., Kneuss, E., Kuncak, V.: Counter-example complete verification for higher-order functions. In: Haller and Miller [6], pp. 18–29. <https://doi.org/10.1145/2774975.2774978>

A Project Setup

In this chapter we look at the practical challenges that could occur by using Stainless. Since Stainless is still in development and there is no 1.x release there might still be breaking changes and improvements fixing problems described now.

nur über stainless-sbt-plugin reden damit man weiss wir haben es versucht

A.1 stainless-sbt-plugin vs JAR

We can either use the sbt plugin or a JAR file to check code with Stainless.

Invoking the JAR on our source code Stainless will verify it. If we have a bigger project, this becomes really tricky, because we must pass all files needed including the dependencies. This is in contrast to the sbt plugin, where we can integrate Stainless in our compilation process. When we call compile, Stainless verifies the code and stops the compilation if the verification fails.

Having a static version configured in the sbt build file, every developer has the same Stainless features available. This should prevent incompatibility with new or deprecated features when we use different plugins.

So the sbt plugin has clear advantages over the JAR file since its integrated directly. We do not have to download it manually and find the right version and if we bump the version we can just edit it in the build file and every developer is on the same version again.

However, currently there are some drawbacks. For example the sbt plugin does not always report errors.

We use the jar for everything.

A.2 Integration into Bitcoin-S

During this work, Stainless updated the sbt plugin to support sbt 1.2.8 from 0.13.17 and Scala 2.12.8 from 2.11.12. So this section might be out of date now.

Stainless requires and Scala recommends Java SE Development Kit 8. Newer Java versions won't work.

To use the latest version of the sbt tool you have to build it locally. You can run `sbt universal:stage` in the cloned Stainless git repository. This generates `frontends/scalac/target/universal/stage/bin/stainless-scalac`.

Bitcoin-S-Core uses sbt 1.2.8 and Scala 2.12.8, while Stainless sbt plugin is on sbt 0.13.17 and Scala 2.11.12.

Sbt introduced new features in the 1.x release used by Bitcoin-S. Most of them can be written the sbt 0.13.17 way.

The bigger problem is, due to the different Scala and sbt versions, the following error after trying to go in a sbt shell:

```
[warn] There may be incompatibilities among your library dependencies; run 'evicted'
       to see detailed eviction warnings.
[error] java.lang.NoClassDefFoundError: sbt/SourcePosition
...
Project loading failed: (r)etry, (q)uit, (l)ast, or (i)gnore?
```

Downgrading Bitcoin-S sbt version to 0.13.17 fixes the error but then it can not load some libraries only compiled for newer versions. So this would take too much time to fix and changes the Bitcoin-S code inadvertently.

The next approach is to use the stainless cli instead of sbt. Running stainless on all source files does not work, because dependencies are missing. The parameter `-classpath` can resolve it but the value of this parameter must be the paths to all the dependencies separated by a `'.'`. Finally, `core` depends on `secp256k1jni`, another package of Bitcoin-S written in Java. So this needs to be in the source files to.

The final command looks like this in `core` folder of Bitcoin-S:

```
$ stainless
  -classpath ".:$(find ~/.ivy2/_-type_f_-name_*.jar | tr '\n' ':') "
  $(find . -type f -name *.scala | tr '\n' ' ')
  $(find ../secp256k1jni -type f -name *.java | tr '\n' ' ')
```

`.ivy2` is the dependency cache of sbt. The `tr` replaces the first char with the second so a newline with either `'.'` or `' '`.

With this command, Stainless throws the next error:

```
[Internal] Error: object scala.reflect.macros.internal.macroImpl in compiler mirror
           not found.. Trace:
[Internal] - scala.reflect.internal.MissingRequirementError$.signal
           (MissingRequirementError.scala:17)
...
[Internal] object scala.reflect.macros.internal.macroImpl in compiler mirror not found.
[Internal] Please inform the authors of Inox about this message
```

So we can not know how many errors will face us. Let's go another way, because the errors may take too much time and it might lead to a next error. We extract the code needed to verify a transaction mainly the class `Transaction` and `ScriptInterpreter` with many other classes they're depending on.

After this extraction Stainless was successfully integrated with both sbt and JAR.

Running `sbt compile` in the project with Stainless ended without error. But it also ended with no output. So we are not able to change the code so Stainless would accept it since we do not know what to change.

So the sbt plugin does not always complain where the JAR file did. The open [issue 484 on GitHub](#) might describe exactly this error.

Now we can finally run Stainless on our code. But this leads us to the next findings. We must rewrite most of the code, as described in the previous chapters.

- clone our repo
- jar vs sbt, you can use both
- contains a stainless jar
- call the script blahblah

B `NumberType.scala`

```

1 package addition.reduced.number
2
3 /**
4  * This abstract class is meant to represent a signed and unsigned
5  *   number in C
6  * This is useful for dealing with codebases/protocols that rely on
7  *   C's
8  * unsigned integer types
9  */
10 sealed abstract class Number[T <: Number[T]] {
11   type A = BigInt
12
13   /** The underlying scala number used to hold the number */
14   protected def underlying: A
15
16   def toLong: Long = toBigInt.bigInteger.longValueExact()
17   def toBigInt: BigInt = underlying
18
19   /**
20    * This is used to determine the valid amount of bytes in a number
21    * for instance a UInt8 has an andMask of 0xff
22    * a UInt32 has an andMask of 0xffffffff
23    */
24   def andMask: BigInt
25
26   /** Factory function to create the underlying T, for instance a
27     UInt32 */

```

```

25   def apply: A => T
26
27   def +(num: T): T = apply(checkResult(underlying + num.underlying))
28
29   /**
30    * Checks if the given result is within the range
31    * of this number type
32    */
33   private def checkResult(result: BigInt): A = {
34     require((result & andMask) == result,
35       "Result was out of bounds, got: " + result)
36     result
37   }
38 }
39
40 /**
41  * Represents a signed number in our number system
42  * Instances of this is [[Int64]]
43  */
44 sealed abstract class SignedNumber[T <: Number[T]] extends Number[T]
45
46 /**
47  * Represents a int64_t in C
48  */
49 sealed abstract class Int64 extends SignedNumber[Int64] {
50   override def apply: A => Int64 = Int64(_)
51   override def andMask = 0xffffffffffffffffL
52 }
53
54 /**
55  * Represents various numbers that should be implemented
56  * inside of any companion object for a number
57  */
58 trait BaseNumbers[T] {
59   def zero: T
60   def one: T
61   def min: T
62   def max: T
63 }
64
65 object Int64 extends BaseNumbers[Int64] {
66   private case class Int64Impl(underlying: BigInt) extends Int64 {
67     require(underlying >= -9223372036854775808L,
68       "Number was too small for a int64, got: " + underlying)
69     require(underlying <= 9223372036854775807L,
70       "Number was too big for a int64, got: " + underlying)
71   }
72
73   lazy val zero = Int64(0)
74   lazy val one = Int64(1)

```

```

75
76 lazy val min = Int64(-9223372036854775808L)
77 lazy val max = Int64(9223372036854775807L)
78
79 def apply(long: Long): Int64 = Int64(BigInt(long))
80
81 def apply(bigInt: BigInt): Int64 = Int64Impl(bigInt)
82 }

```

C CurrencyUnits.scala

```

1 package addition.reduced.currency
2
3 import addition.reduced.number.{BaseNumbers, Int64}
4
5 sealed abstract class CurrencyUnit {
6   type A
7
8   def satoshis: Satoshi
9
10  def !=(c: CurrencyUnit): Boolean = !(this == c)
11
12  def ==(c: CurrencyUnit): Boolean = satoshis == c.satoshis
13
14  def +(c: CurrencyUnit): CurrencyUnit = {
15    Satoshi(satoshis.underlying + c.satoshis.underlying)
16  }
17
18  protected def underlying: A
19 }
20
21 sealed abstract class Satoshi extends CurrencyUnit {
22   override type A = Int64
23
24   override def satoshis: Satoshi = this
25
26   def toBigInt: BigInt = BigInt(toLong)
27
28   def toLong: Long = underlying.toLong
29
30   def ==(satoshis: Satoshi): Boolean = underlying == satoshis.
    underlying
31 }
32
33 object Satoshi extends BaseNumbers[Satoshi] {
34
35   val min = Satoshi(Int64.min)
36   val max = Satoshi(Int64.max)
37   val zero = Satoshi(Int64.zero)
38   val one = Satoshi(Int64.one)

```

```

39
40   def apply(int64: Int64): Satoshi = SatoshiImpl(int64)
41
42   private case class SatoshiImpl(underlying: Int64) extends Satoshi
43 }

```

D Code Transformations

Here we see in detail how to transform the bitcoin-s code into the Scala fragment supported by Stainless. All subsections start with the Stainless error message(s) and finally a description of the changes we make to the code.

We claim that all transformations are equivalent in the sense that if the addition-with-zero property holds for the transformed code, then it also holds for the code before the transformation.

D.1 Inheriting Objects

```

[ Error ] number/NumberType.scala:65:1: Objects cannot extend
        classes or implement traits, use a case object instead
        object Int64 extends BaseNumbers[Int64] {
          .....
[ Error ] currency/CurrencyUnits.scala:33:1: Objects cannot extend
        classes or implement traits, use a case object instead
        object Satoshi extends BaseNumbers[Satoshi] {
          .....

```

Here, we can just turn the objects into case objects by literally just changing the word `object` into `case object` on lines 65 and 33 in the two respective files.

That transformation is equivalent. Case objects have some additional properties (in particular, being serializable) and they inherit from `Product` instead of `AnyRef`, but none of our code depends on any of that.

D.2 Abstract Type Members

```

[ Error ] currency/CurrencyUnits.scala:6:3: Stainless doesn't
        support abstract type members
        type A
        .....

```

Note that we can not simply replace the unsupported abstract type member by a (supported) type parameter. The problem is that the `CurrencyUnit` class uses one of its implementing classes: `Satoshi`.

`Satoshi` would have to instantiate a potential type parameter with type `Int64`, so it would extend `CurrencyUnit[Int64]`. But that is too specific, because the return type of the `+`-method would then be `CurrencyUnit[Int64]` not `CurrencyUnit[A]`.

Since we only want to verify the addition of `satoshi`, and the `Satoshi` class overrides `A` with `Int64` anyway, we just remove the abstract type and set it to `Int64`.

We remove line 6 and line 22 from `CurrencyUnits.scala` (to maintain line numbers we actually replace them with empty lines for now) and in line 18 we replace `A` by `Int64`.

D.3 Non-Literal BigInt Constructor Argument

```
[ Error ] currency/CurrencyUnits.scala:26:33: Only literal arguments
         are allowed for BigInt.
         def toBigInt: BigInt = BigInt(toLong)
                                   ^^^^^^^
```

As described before, the types in the Stainless library are more restricted than their Scala library counterparts. In particular, the Stainless `BigInt` constructor is restricted to literal arguments.

After:

```
1 sealed abstract class Satoshi extends CurrencyUnit {
2   override def satoshis: Satoshi = this
3
4   def toBigInt: BigInt = underlying.toBigInt
5
6   def toLong: Long = underlying.toLong
7
8   def ==(satoshis: Satoshi): Boolean = underlying == satoshis.
      underlying
9 }
```

Here we can simply use `toBigInt` on the field `underlying` directly. So, instead of converting the underlying to `Long` and back to `BigInt` we convert underlying directly to `BigInt`.

We replace line 26 by

```
1 def toBigInt: BigInt = underlying.toBigInt
```

This is an equivalent transformation: the only thing that might go wrong in the detour via `Long` is that the `BigInt` underlying `Int64` in turn underlying `Satoshi` does not fit into a `Long`. However, the only constructor of `Int64` ensures exactly that.

D.4 Self-Reference in Type Parameter Bound

```
[ Error ] number/NumberType.scala:8:30: Unknown type parameter type
      T
      sealed abstract class Number[T <: Number[T]] {
                                   ^^^^^^^^^^^^^^^^^
```

Stainless does not currently support a class with a type parameter with a type boundary that contains the type parameter itself. We opened an issue [\[3\]](#) on the Stainless repository and the Stainless developers have targeted Stainless version 0.4 to support such self-referential type boundaries.

For now, since our code only uses `Number` with type parameter `T` instantiated to `Int64`, we just remove the type parameter and replace it by `Int64`. We respectively replace lines 8, 44 and 49 by

```

1 sealed abstract class Number {

1 sealed abstract class SignedNumber extends Number

1 sealed abstract class Int64 extends SignedNumber {

```

and replace T by Int64 in lines 25-27.

D.5 Missing Member `bigInteger` in `BigInt`

```

[ Error ] number/NumberType.scala:14:22: Unknown call to bigInteger
        on Number.this.toBigInt (BigInt) with arguments List()
        of type List()
        def toLong: Long = toBigInt.bigInteger.longValueExact()
                        ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

```

stainless error emssage missing

The Scala class `BigInt` is essentially a wrapper around `java.math.BigInteger`. `BigInt` has a member `bigInteger` which is the underlying instance of the Java class. The Java class has a method `longValueExact` which returns a long only if the `BigInteger` fits into a long, otherwise throws exception. Stainless does not support Java classes and in particular its `BigInt` has no member `bigInteger`.

However, our code never calls `toLong` anymore, so we just remove it. We replace line 14 in `NumberType.scala` and line 28 in `CurrencyUnits.scala` by an empty line.

D.6 Type Member

```

[Warning ] number/NumberType.scala:9:3: Could not extract tree in
        class: type A = BigInt (class scala.reflect.internal.
        Trees$TypeDef)
        type A = BigInt
        ^^^^^^^^^^^^^^^^^

```

Our version of Stainless does not support type members. We just replace all occurrence of `A` with `BigInt`, since `A` is never overwritten in an implementing class.

We remove line 9 in `NumberType.scala` and replace `A` by `BigInt` in lines 12, 25, 33 and 50.

In the mean time Stainless has implemented support for type member [4]. Since version 0.2 verification should succeed without this change.

D.7 Missing Bitwise-And Method on `BigInt`

```

[ Error ] number/NumberType.scala:34:14: Unknown call to & on result
        (BigInt) with arguments List(Number.this.andMask) of
        type List(BigInt)
        require((result & andMask) == result,
                ^^^^^^^^^^^^^^^^^^^^^^^^^

```

Contrary to Scala BigInt, the Stainless BigInt class does not support bitwise operations, in particular not the &-method for bitwise and.

The bitwise and with the andMask on line 34 is a bounds check. It checks if the result parameter is in range of the specified type, which in our case is the hard coded Int64.

So, we replace the `&` mask with a check whether the result is in range of `Long.MinValue` and `Long.MaxValue`, because `Int64` has the same 64-bit range as `Long`.

We replace lines 34-35 by the following, squeezed into two lines to preserve line numbers. Note that the BigInt constructor requires a literal:

```
1 require(result <= BigInt("9223372036854775807")
2    && result >= BigInt("-9223372036854775808"),
3    "Result was out of bounds, got: " + result)
```

kai here we do not know if it changes semantic. ramon: oh, we don't?

D.8 Inner Class in Case Object

[illegible]

Stainless does not support inner classes in a case object. Bitcoin-s uses this a lot to separate the class

ramon: i don't understand from its implementation.

This is easy to fix. We just move the inner classes out of the case objects. They do not interfere with any other code.

We remove lines 66-71 in `NumberType.scala` and insert them at the end of the file. We remove line 42 in `CurrencyUnits.scala` and insert it at the end of the file.

D.9 Message Parameter in Require

Here the `require(condition, message)` is used as an assertion: if the condition is false the fail with message. Stainless does not support the message parameter of `require`. For the verification we can simply remove that parameter.

```
[ Error ] inv$4 depends on missing dependencies: long2bigInt$0.
```

We replace the two `require` clauses at lines 67 and following in `NumberType.scala`

We also replace lines 76 and 77

D.11 Missing BigInt Constructor with Long Argument

```
[ Error ] apply$14 depends on missing dependencies: BigInt$0,
        apply$15.
```

Here again, the Scala `BigInt` has a constructor with a `Long` argument which is missing in the `Stainless BigInt`.

We simply replace the `Long` values in a `BigInt` constructor call with a string literal.

The lines from the previous transformation respectively change into the following:

```
1   require(underlying >= BigInt("-9223372036854775808"))
2   require(underlying <= BigInt("9223372036854775807"))
   and
1   lazy val min = Int64(BigInt("-9223372036854775808"))
2   lazy val max = Int64(BigInt("9223372036854775807"))
3 }
```